TRAINING DECISION

Untitled design

Prepared 11 Oct 2026, 18:39 · figures from the analysis of 11 Oct 2026, 18:38 · proposal p-7c0cebc6e85e2125

RECOMMENDED — READY TO APPROVE

$0.00 to train, over 13 min.

- Fits the chosen target with 5.5 GB to spare.
- NEURAX proposes a budget of $10 — the top of the $0.00–$0.00 range, plus 15% for runs that fail and restart.

BUDGET TO APPROVE	ESTIMATED COST	DURATION	HARDWARE
$10	**$0.00**	**13 min**	**Intel(R) Core(TM) i5-8365U CPU @ 1.60GHz**

THE MODEL

Name	**Untitled design**	
Architecture	**Transformer**	
Trained to	**Regression**	As the output block states it.
Parameters	**43.4 k**	Exact: counted by NEURAX from the design.
Weights	**170 kB**	At fp32.

THE DATA

Dataset	**sample_10k_rows.csv**	On this machine.
Kind	**CSV table**	
Samples	**10,000**	
Size	**14 MB**	On disk. A run reads it from where it is; it is not copied.
One sample	**273**	
Target	**wealth_building_index**	The column predicted, as the profiler read it.
Integrity	**17,764 missing values**	17764 missing values in the first 5000 rows.
Split	**9,000 train · 1,000 validation**	The validation part is held out and evaluated every epoch.
Checked	**Read from the data itself**	Fingerprint 79414208e8b4.

THE TRAINING RUN

Length	**20 epochs × 282 steps = 5,640 steps**	Over 9,000 training samples.
Batch	**32 per step**	
Input	**273 features**	One sample, as the model reads it.
Optimizer	**adamw · learning rate 0.001**	Weight decay 0.1, 564 warmup steps, cosine schedule.
Precision	**fp32**	
Data split	**90% train · 10% validation · 0% test**	The runtime records the effective sample counts and evaluates the held-out test part after training.
Memory saving	**None**	
Early stopping	**Off — the run takes its full length**	

Seed	**0**	The same seed, data and settings repeat the run.
Checkpoints	**Every 564 steps**	About 509 kB each.

WHERE IT RUNS

Hardware	**Intel(R) Core(TM) i5-8365U CPU @ 1.60GHz**	The machine the analysis was computed for.
GPUs	**1**	
Location	**This machine**	
Price	**—**	No rented instance chosen.
Speed	**Calibrated on 1 run of this machine**	calibrated on 1 run of this machine: real steps took 4.36× the prediction (±10%)
Peak memory	**0.0 GB**	Against 5.5 GB available.

WHAT IT COSTS

Training cost	**$0.00**	Between $0.00 and $0.00. 13 min on Intel(R) Core(TM) i5-8365U CPU @ 1.60GHz
Budget proposed	**$10**	$0.00 if the steps run at the slow end of the range, plus 15% for runs that fail and restart. The run goes to its end either way.
Duration	**13 min**	Between 12 min and 14 min. Wall-clock, at the predicted step time.
Energy	**0.0 kWh**	The device power draw over the predicted duration.
Carbon	**0.0 kg CO2e**	The energy above at the grid intensity given.

READY TO START ONCE APPROVED

Figures are current	**Yes**	Every figure is the analysis of the design on screen.
NEURAX can analyse the design	**Yes**	No unresolved error.
It trains on the machine it was priced for	**Yes**	Priced for this machine, where it trains.
The model predicts every class of the data	**Yes**	Its output covers the classes of the data.
It fits the target's memory	**Yes**	5.5 GB to spare.
Training data is chosen, on this machine	**Yes**	sample_10k_rows.csv.
The run has a folder	**Yes**	Written into runs/ in /home/fossouomartial/NEURAX/Untitled design, chosen by NEURAX. Save the design to choose another folder.

RISKS AND OPEN POINTS

WARNING · The data has problems

The profiler found 17,764 missing values in sample_10k_rows.csv.

Fix or remove them before training: a run skips nothing it cannot read without saying so.

HINT · Memory and time are estimates

Parameters and FLOPs are exact. Kept activations follow each operation's backward rules. Step time uses the machine's throughput and is corrected only by completed runs at a comparable predicted scale.

A training run measures both, and Training -> Accuracy compares them with these figures.

WHAT THESE FIGURES ARE

• Parameters and FLOPs are exact. Time, energy and money are estimated from the machine’s measured or published throughput, and each figure says which.

• Time, energy and money are given as a range: a step takes what training really reaches of the machine — 25–55% of a published peak, 50–90% of a rate measured here — until finished runs on this machine calibrate it. The proposed budget sits at the top of that range, plus a stated contingency for runs that fail and restart — 15% for a short run, rising to 128% at frontier scale — a convention, not a measurement of this design.

• A run measures its real step time, memory and cost; Training compares them with these figures.

• Energy and carbon are per training run, at the device’s power draw and the grid intensity given. Your supplier’s intensity and your data centre’s PUE change the figure a CSRD disclosure uses.

• EU AI Act classification follows from the use of the model — who is affected and which decision its output feeds — not from its architecture. It belongs with the people who know the deployment.

DECISION

Outcome	**Approved**	
By	**martial wato**	11 Oct 2026, 18:39
Reason	**go**	